




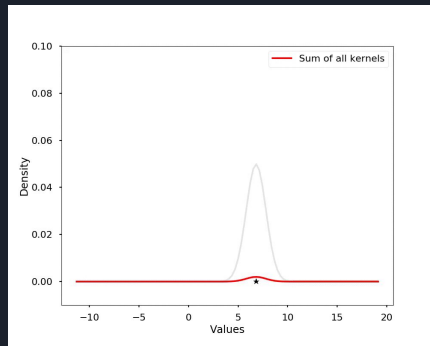
Popping the Kernel

Using Constrained Optimization to Estimate
the Values behind a Kernel Estimation

By Aneesh Sandhir

What is Kernel Estimation?

- Kernel Estimation is a transformation of a dataset through the application of a distribution, the kernel
- Kernels are defined by their shape which determines the properties of the resulting transform
- A kernel's width is its most important attribute
- Kernels can be useful tools to
 - Investigate spatial or cyclic patterns
 - Remove noise
 - Develop a probability density function





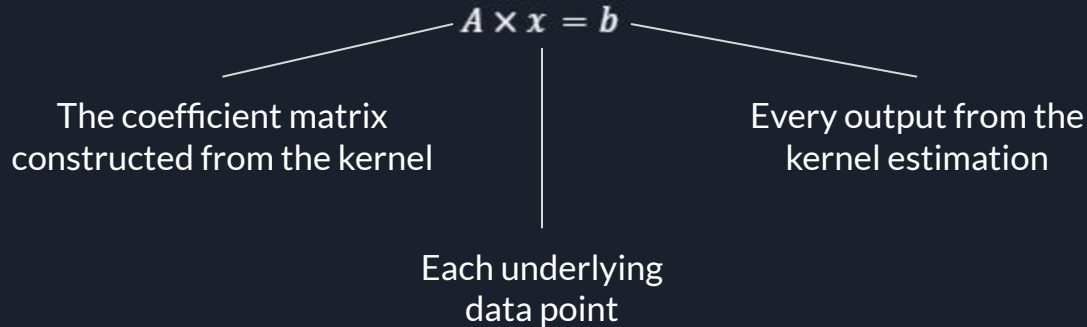
Why Undo a Kernel Estimation?

- The kernel obfuscates the data
- The kernel applied to the data can constrain your analysis
- There is almost no way to apply a different kernel to the data without the raw data itself
 - Only if the new kernel's size is a multiple of the kernel already applied to the data, this can be very limiting




How to Undo a Kernel Estimation?


- Each output of a Kernel Estimation can be viewed as its own linear equation
- And with many of the linear equations sharing variables, the entire transformation can be viewed as the following system of linear equations



How to Undo a Kernel Estimation? Continued

- Due to the nature of kernel density estimations, the system will have more variables than constraining equations, this means there are infinitely many solutions to the system
- This solution space can be reigned in by further constraining the system
- These constraints can take the form of inequalities, setting upper and lower bounds on the variables, or number classification, either continuous, integer, or binary





Where Constrained Optimization Fits into the Data Engineer's Toolbox?

Data Transformation	Action	Reaction
Row Wise	Union	Filter
Column Wise	Join	Subset
Reshaping	Pivot	Melt
Accumulation	Cumulative Sum	Lag
Rolling Aggregations	Kernel Estimation	Constrained Optimization



One Dimensional Example: Time Series - CDC Overdose Data



Vital Statistics Rapid Release Drug Overdose Data

- The VSRR data tracks the mortality statistics indexed by region, month and drug, spanning from 2015 to 2022
- The mortality statistics are recorded as the sum of the current month's death count and the death count of the preceding 11 months
- This hides some of the month-to-month variation and makes it easier to identify a general trend but obfuscates the underlying data and constrains the analysis through a 12-month window
- Each month's exact death count can't be determined, they can be estimated through constrained optimization



Developing the System of Equations

- A 12-month rolling sum is simply a uniform kernel with a width of 12 and a height of 1
- Developing the system starts by representing a given month as i , its 12-month rolling death tally Σ_i , the death tally for that month x_i , and the death tally for the preceding 11 months as $x_{i-11} \dots x_i$ will yield the following constraining equation


$$x_{i-11} + x_{i-10} + x_{i-9} + x_{i-8} + x_{i-7} + x_{i-6} + x_{i-5} + x_{i-4} + x_{i-3} + x_{i-2} + x_{i-1} + x_i = \Sigma_i$$

- The next month can be represented as $i+1$, its 12-month rolling death tally as Σ_{i+1} , the death tally for that month as x_{i+1} and will yield the following constraining equation

$$x_{i-10} + x_{i-9} + x_{i-8} + x_{i-7} + x_{i-6} + x_{i-5} + x_{i-4} + x_{i-3} + x_{i-2} + x_{i-1} + x_i + x_{i+1} = \Sigma_{i+1}$$

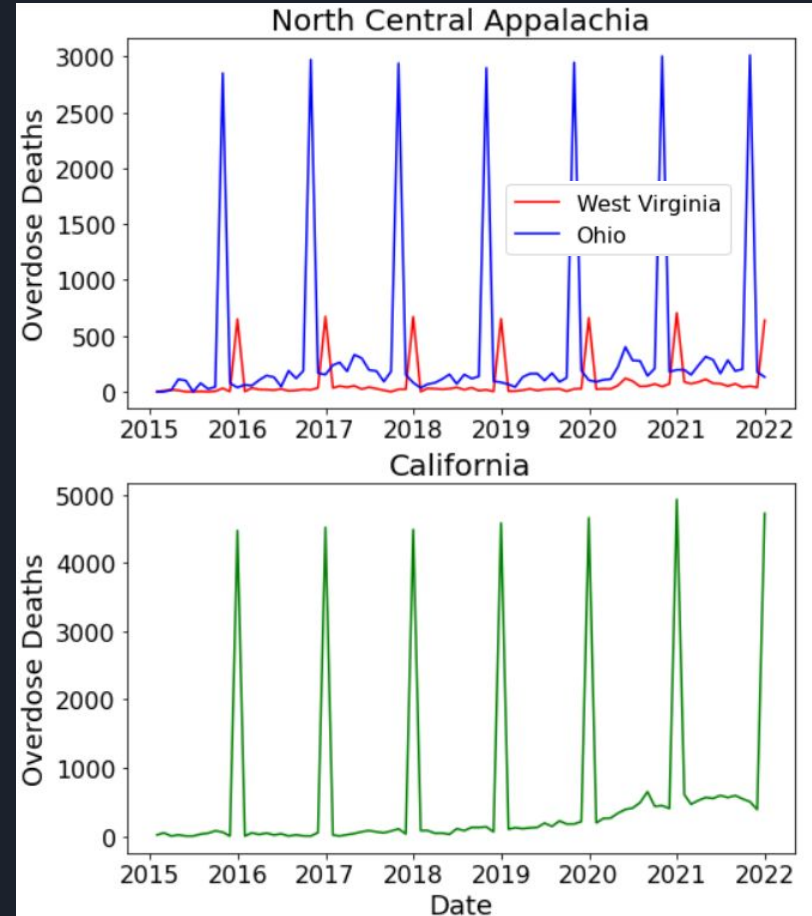
Constraining the System

- The system has 11 more variables than constraining equations
- Knowing the variables represent overdose deaths the solution space can be constrained to a finite set
- These additional constraints can be implemented by first confining the solution space to integers
- Followed by applying a constraining inequality where none of the elements in the variable matrix can be less than 0




Initial Solution is Unrealistic

- Solving this system via optimization yields the least entropic, most volatile solution with an overwhelming majority of deaths occurring one month a year
- In all likelihood, the system does not behave this way




Further Constraining the System

- To get a less volatile solution, the system can be further constrained by applying another constraining inequality where none of the elements in the variable matrix can be exceed a given value
- Here it is implemented as a maximum percentage deviation **above** the 12-month moving average, also referred to as a smoothing or volatility factor



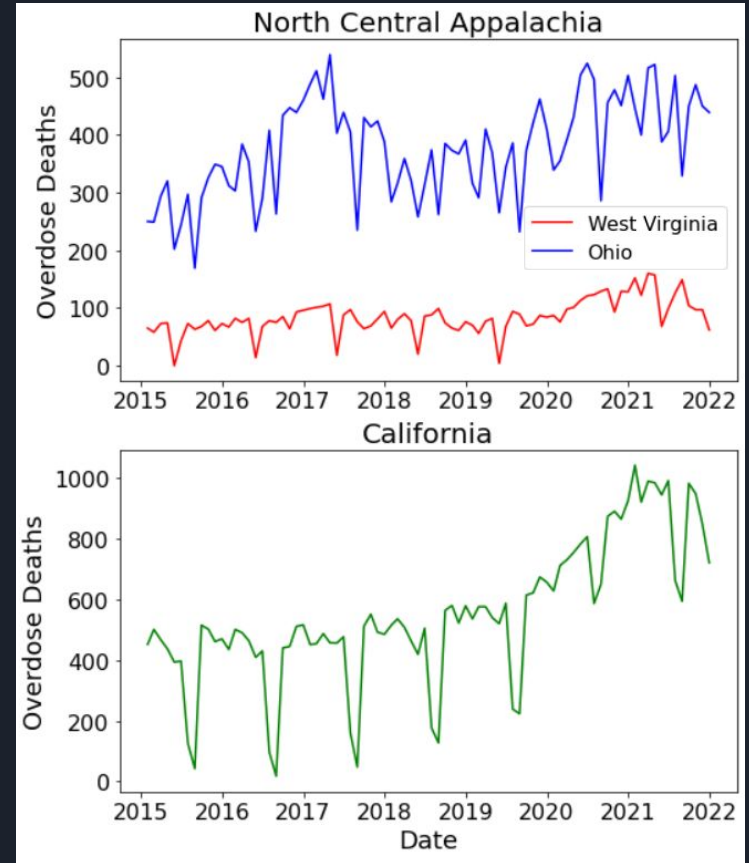
Low Volatility Factor - Overconstrained System

- If the volatility factor is too low, tightly tying the monthly values to the moving average, the system will be overly constrained and no solution can satisfy all the constraints put on the system



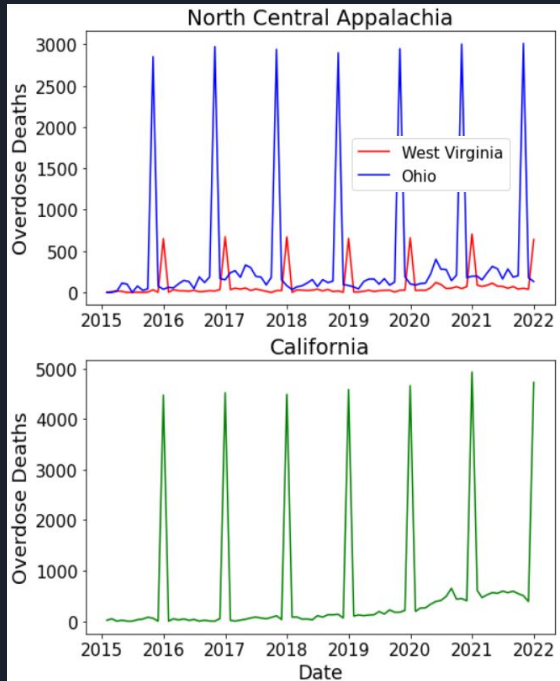
Optimal Volatility Factor

- If the volatility factor is too high the resulting solution will be artificially cyclical and very volatile, as was previously
- Selecting the right volatility factor depends on the ratio of unknown values to the constraining equations, the magnitude of the values in the series and can only be identified via trial and error

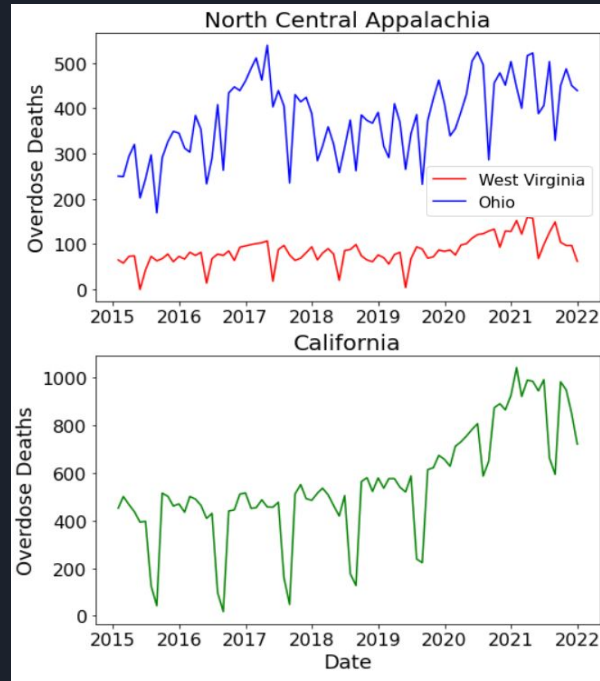


Comparison of Constraints

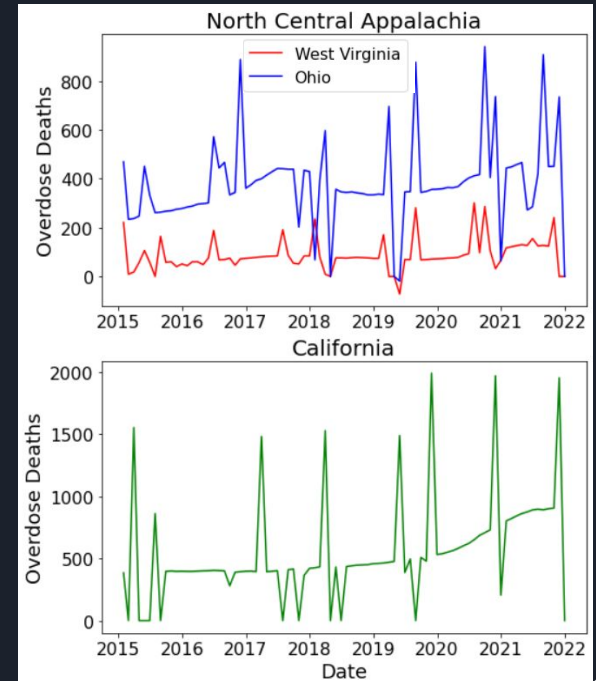
Under Constrained



Optimal Volatility Factor




Overly Constrained




Additional Constraints - Dynamic Lower Bound

- Applying a dynamic upper bound to the variables results in a system that behaves a lot more realistic, but there are still drastic sharp rises and drops that appear at almost the same time every year
 - Instead of applying a static lower bound of 0 to all the variables, the lower bound can be made dynamic as well
- Using the same volatility factor as the upper bound, the variables can have a maximum percentage deviation **below** the 12-month moving average

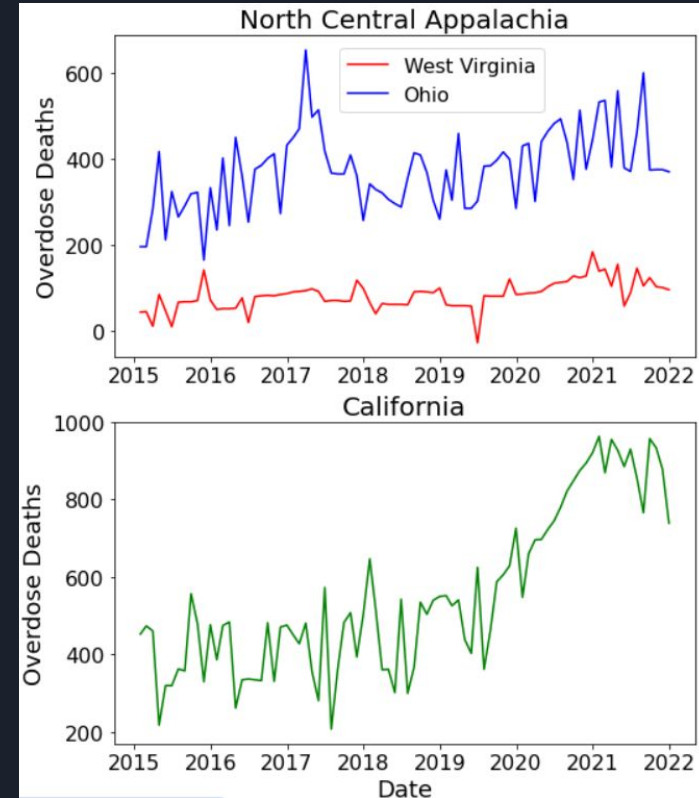


Comparison of Constraints on the Lower Bound

Non-Negative Lower Bound

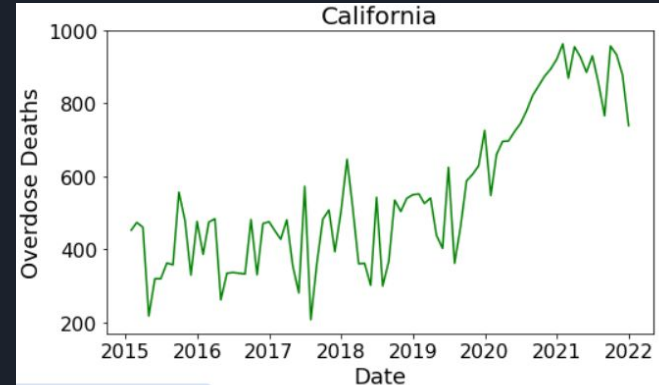
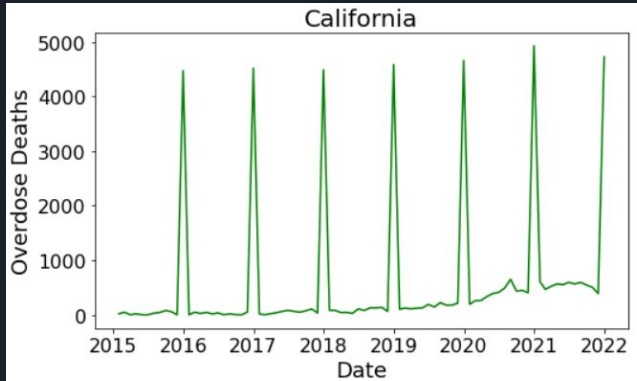


Dynamically Constrained Lower Bound



Conclusions

- By fine tuning the constraints applied to the system and identifying optimal volatility factor, we found a good estimate of the underlying data



- Transformed 12-month rolling sum into month by month data
- This monthly overdose data was used as as input into a descriptive model
- Models that used monthly overdose data outperformed the models which used rolling sum data

Initial Python Implementation

```
1
2 #Import dependencies
3 import pandas as pd
4 import numpy as np
5 from datetime import datetime, timedelta
6 import pulp
7 import time
8 import matplotlib.pyplot as plt
9
10 #Import VSRR data
11 ...
12
13 #Create sliding window function
14 def sliding_windows(kernel, num_constraining_eqs):
15     kernel = np.asarray(kernel)
16     p = np.zeros(num_constraining_eqs-1, dtype=kernel.dtype)
17     b = np.concatenate((p, kernel, p))
18     s = b.strides[0]
19     strided = np.lib.stride_tricks.as_strided
20     return strided(b[num_constraining_eqs-1:], shape=(num_constraining_eqs, len(kernel)+num_constraining_eqs-1), strides=(-s, s)
21
```

Initial Python Implementation Continued

```
54 period = 12
55 regions = death_df["Region"].unique()
56 indicies = ["State", "FIPS", "Region", "Year", "Month", "Start Date", "End Date", "Percent with drugs specified"]
57 targets = [element for element in death_df.columns.tolist() if element not in indicies]
58
59 monthly_deaths_df = pd.DataFrame()
60 for region in regions:
61     series = death_df[death_df["Region"] == region].sort_values(by='End Date')
62     for target in targets:
63         sol = series[target].dropna().to_numpy()
64         if (len(sol) < 2):
65             marginal_values = np.empty(len(series.index))
66             marginal_values[:] = np.nan
67             series["Estimated Monthly Marginal " + target] = marginal_values
68             continue
69         coef = sliding_windows(np.ones(period), len(sol))
70         #pulp Linear Algebra Solver yeilds a solution not guaranteed to be the correct solution but all the values
71         #it produces will be non-negative, discrete and yeild the values in the dataset when summed across a 12 month window
72         mod = pulp.LpProblem(region.replace(" ", "_") + "." + target.replace(" ", "_"))
73
74         #set up contraining inequalities
75         #constrain monthly death tally values to integers greater than or equal to 0
76         vars = pulp.LpVariable.dicts('x', range(len(coef[0])), lowBound=0, cat='Integer')
77
78         #set up contraining equations
79         for row, rhs in zip(coef, sol):
80             mod += sum([row[i]*vars[i] for i in range(len(row))]) == rhs
81
82         mod.solve()
83
84         #currently discarding the leading 11 values which dont have corresponding 12 month rolling sums
85         #could possibly rewrite this section to keep them with null values for the 12 month rolling sums in the future
86         marginal_values = [vars[i].value() for i in range(len(coef[0]))][(period -1):]
87         series["Estimated Monthly Marginal " + target] = np.pad(marginal_values, (len(series.index)-len(marginal_values),0),
88
89         monthly_deaths_df = pd.concat([monthly_deaths_df, series])
```

Python Implementation of More Constraints

```
120 period = 12
121 volatility_factor = 1.2 # must be between 1 and period
122 regions = death_df["Region"].unique()
123 indices = ["State", "FIPS", "Region", "Year", "Month", "Start Date", "End Date", "Percent with drugs specified"]
124 targets = [element for element in death_df.columns.tolist() if element not in indices]
125
126 monthly_deaths_df = pd.DataFrame()
127 for region in regions:
128     series = death_df[death_df["Region"] == region].sort_values(by='End Date')
129     for target in targets:
130         sol = series[target].dropna().to_numpy()
131         if (len(sol) < 2):
132             marginal_values = np.empty(len(series.index))
133             marginal_values[:] = np.nan
134             series["Estimated Monthly Marginal " + target] = marginal_values
135             continue
136         coef = sliding_windows(np.ones(period), len(sol))
137         #pulp Linear Algebra Solver yeilds a solution not guaranteed to be the correct solution but all the values
138         #it produces will be non-negative, discrete and yeild the values in the dataset when summed across a 12 month window
139         mod = pulp.LpProblem(region.replace(" ", "_") + ":" + target.replace(" ", "_"))
140
141         #set up contraining inequalities
142         #constrain monthly death tally values to integers greater than or equal to 0
143         vars = pulp.LpVariable.dicts('x', range(len(coef[0])), lowBound=0, cat='Integer')
144
145         #constrain monthly death tally values to be less than n% greater than themoving average
146         upBound_vector = np.around(((volatility_factor/period)*np.pad(sol, (period-1,0), 'constant', constant_values=sol[0]))
147         for pointer in vars.keys():
148             vars[pointer].upBound = upBound_vector[pointer]
149
150         #set up contraining equations
151         for row, rhs in zip(coef, sol):
152             mod += sum([row[i]*vars[i] for i in range(len(row))]) == rhs
153
154         mod.solve()
```

Python Implementation of More Refined Constraints

```
182 volatility_factor = 1.2 # must be between 1 and period
183 regions = death_df["Region"].unique()
184 indices = ["State", "FIPS", "Region", "Year", "Month", "Start Date", "End Date", "Percent with drugs specified"]
185 targets = [element for element in death_df.columns.tolist() if element not in indices]
186
187 monthly_deaths_df = pd.DataFrame()
188 for region in regions:
189     series = death_df[death_df["Region"] == region].sort_values(by='End Date')
190     for target in targets:
191         sol = series[target].dropna().to_numpy()
192         if (len(sol) < 2):
193             marginal_values = np.empty(len(series.index))
194             marginal_values[:] = np.nan
195             series["Estimated Monthly Marginal " + target] = marginal_values
196             continue
197         coef = sliding_windows(np.ones(period), len(sol))
198         #pulp Linear Algebra Solver yeilds a solution not guaranteed to be the correct solution but all the values
199         #it produces will be non-negative, discrete and yeild the values in the dataset when summed across a 12 month window
200         mod = pulp.LpProblem(region.replace(" ", "_") + ":" + target.replace(" ", "_"))
201
202         #set up contraining inequalities
203         #constrain monthly death tally values to integers greater than or equal to 0
204         vars = pulp.LpVariable.dicts('x', range(len(coef[0])), lowBound=0, cat='Integer')
205
206         #constrain monthly death tally values to be less than n% greater than themoving average
207         lowBound_vector = np.around(((1/(volatility_factor*period))*np.pad(sol, (period-1,0), 'constant', constant_values=sol
208         upBound_vector = np.around(((volatility_factor/period)*np.pad(sol, (period-1,0), 'constant', constant_values=sol[0]))
209         for pointer in vars.keys():
210             vars[pointer].lowBound = lowBound_vector[pointer]
211             vars[pointer].upBound = upBound_vector[pointer]
212
213         #set up contraining equations
214         for row, rhs in zip(coef, sol):
215             mod += sum([row[i]*vars[i] for i in range(len(row))]) == rhs
216
217         mod.solve()
218
```




Developing the System of Equations

- A 12-month rolling sum is simply a uniform kernel with a width of 12 and a height of 1
- Developing the system starts by representing a given month as i , its 12-month rolling death tally Σ_i , the death tally for that month x_i , and the death tally for the preceding 11 months as $x_{i-11} \dots x_i$ will yield the following constraining equation

$$x_{i-11} + x_{i-10} + x_{i-9} + x_{i-8} + x_{i-7} + x_{i-6} + x_{i-5} + x_{i-4} + x_{i-3} + x_{i-2} + x_{i-1} + x_i = \Sigma_i$$

- The next month can be represented as $i+1$, its 12-month rolling death tally as Σ_{i+1} , the death tally for that month as x_{i+1} and will yield the following constraining equation

$$x_{i-10} + x_{i-9} + x_{i-8} + x_{i-7} + x_{i-6} + x_{i-5} + x_{i-4} + x_{i-3} + x_{i-2} + x_{i-1} + x_i + x_{i+1} = \Sigma_{i+1}$$